

Subword Tokenizer & Prototype Training

Production-validated 35.2M-param TinyLM (0.0107 eval loss)

July 8, 2026

1. Executive Summary

Project: Production-validated Subword Tokenizer + SLM Prototype Training and Evaluation

Status: Phase 3 Complete (July 13, 2026)

Key Results:

- **Phase 1**: 0.0107 eval loss on 6000-step TinyLM (35.2M params)
- **Phase 2**: best LoRA eval loss **0.0060** (44% improvement vs Phase 1)
- **Phase 3**: full 7-task evaluation suite complete (benchmark, qualitative, test set, PPL/BLEU, physics QA)

System Architecture

Two unified layers:

1. **Tokenizer layer** (Rust, production-grade)
 - BPE tokenizer in `subword_tokenizer/` with 32K vocabulary
 - CLI + library API (train/encode/decode/model operations)
 - Tested with physics corpus and arXiv papers
2. **Training layer** (Python, production-validated)
 - `stream_train.py`: PyTorch TinyLM with live arXiv streaming
 - MPS acceleration on Apple Silicon (M3 Pro tested)
 - Network resilience: 100-retry exponential backoff
 - Checkpoint periodicity: 500-step saves + final checkpoint

Validation Results (Phase 1-3)

Production Checkpoint: `prototype_long4h_epoch1.pt`

- **Model**: TinyLM (d_model=384, n_layers=6, n_heads=6, 35.2M params)
- **Training**: 6000 steps on physics arXiv corpus
- **Eval Loss**: 0.0107 (50-batch validation, 256-token sequences)
- **Runtime**: 3.5 hours on Apple M3 Pro (2 sec/step)
- **Stability**: Zero NaN, zero OOM, zero crashes
- **Improvement**: 473× better than 2-step baseline (5.0738 → 0.0107)

Phase 2 + 3 validated outcomes (July 13, 2026):

- Test-set loss: 0.010910 → **0.005952** (**45.45%** gain)
- Inference latency benchmark (50 prompts): **1.1449x** average speedup
- Physics QA (20 questions, ranking rubric): 0.325 → **0.450** (Δ **+0.125**)
- BLEU-4 and open-ended qualitative scoring currently constrained by early-EOS generation behavior

Status: **Ready for Phase 4** (RAG integration + generation control tuning)

2. Recent Milestones

July 13, 2026: Phase 3 Completion

- Implemented and validated Tasks 1-7 end-to-end
- Added inference/benchmark/qualitative/test-set/language-metrics/QA scripts

- Produced consolidated report: `PHASE\_3\_COMPLETION.md`
- Confirmed strong retained LoRA gains on held-out test split and QA scoring

****July 8, 2026****: Big Prototype Training & Validation

- Completed 6000-step training run (3.5 hours on M3 Pro)
- Evaluated checkpoint against physics corpus (50 validation batches)
- Achieved 0.0107 eval loss with stable convergence
- Baseline comparison: 473× improvement vs 2-step smoke test
- Committed results to GitHub (commit 4d3f0eb)

****July 2-3, 2026****: Network Resilience & Evaluation Automation

- Implemented exponential backoff retry logic (100 max retries, 5s base)
- Added checkpoint evaluation pipeline (`evaluate\_checkpoints.py`)
- Added automated 3x-run loop with best checkpoint selection
- Validated training stability over 3+ hour unattended runs

****June 30, 2026****: PyTorch + MPS Integration

- Implemented streaming trainer with live arXiv integration
- Added MPS device support for Apple Silicon acceleration
- Verified model loads and trains without OOM on M3 Pro (18GB RAM)
- Created laptop-scale preset configurations (batch_size=4, seq_len=256)

3. Current Architecture

Data Pipeline

1. ****Live Source****: arXiv physics papers (2024 abstracts + intros)
2. ****Retry Logic****: 100-retry exponential backoff (5s base delay)
3. ****Tokenization****: Rust BPE tokenizer (model_32k.json, 32K vocab)
4. ****Training****: PyTorch streaming trainer (8 batches/chunk, dynamic padding)
5. ****Checkpointing****: Step saves every 500 steps + final epoch checkpoint

Model Configuration (Production)

...

TinyLM

```

■■■ d_model: 384
■■■ n_layers: 6
■■■ n_heads: 6
■■■ vocab_size: 32000
■■■ total_params: 35.2M
■■■ batch_size: 4
■■■ seq_len: 256
■■■ learning_rate: 2e-4 (cosine schedule)
■■■ training_steps: 6000

```

...

Device Support

- ****Primary****: Apple M3 Pro (MPS device, 18GB unified memory)

- **Fallback**: CUDA (NVIDIA GPUs) or CPU
- **Status**: M3 Pro fully validated (3.5h runtime, zero OOM)

Architecture diagram: See [doc/architecture.md](architecture.md)

- Multi-run ranking by eval loss (`best\_checkpoint\_\*.json`).

4. Tokenizer CLI Commands (Current)

From `subword\_tokenizer/`:

- `cargo run --bin bpe-tokenizer -- train`
- `cargo run --bin bpe-tokenizer -- expand`
- `cargo run --bin bpe-tokenizer -- tokenize`
- `cargo run --bin bpe-tokenizer -- decode`
- `cargo run --bin bpe-tokenizer -- count`
- `cargo run --bin bpe-tokenizer -- prepare --train --val --test`

Notes:

- Active runtime model file for CLI is `model.json`.
- Training scripts activate the desired model by copying from `model\_32k.json` to `model.json` before tokenization.

5. Prototype Training Commands (Current)

From repository root:

- Quick prototype: `./scripts/run\_prototype.sh`
- 3-run compare and auto-select: `./scripts/run\_prototype\_3x\_and\_select\_best.sh`
- Long run (~3–4h target on M3 Pro): `./scripts/run\_prototype\_long\_4h.sh`
- Evaluate existing checkpoints: `./scripts/run\_evaluation.sh`

6. Validation Status

- Rust tokenizer tests are passing (`src/tests.rs`).
- Prototype training on Mac MPS is validated.
- Best-checkpoint reports are being generated under `checkpoints/`.
- Long-run path now includes network timeout retry/backoff and step checkpoint safety.

7. Current Risks / Observations

1. **Live-stream dependency risk**: arXiv network timeouts can still pause long runs (now retried).
2. **Artifact growth**: checkpoints and corpora can grow quickly.
3. **Documentation drift risk**: ensure markdown-to-PDF sync is part of workflow.

8. Recommended Next Actions

1. Add offline dataset training mode (train from local predownloaded corpus).

2. Add resume-from-checkpoint support for interrupted long runs.
3. Add CI/doc task to regenerate this PDF from markdown automatically.
4. Track large artifacts through retention/LFS policy.

9. Project Roadmap: Physics Research Assistant SLM

9.1 Final Goal

Build a Physics Research Assistant using a 3B-7B parameter SLM, fine-tuned with LoRA on physics papers, combined with RAG and tool-using agents.

9.2 Current Status

Completed:

- Subword tokenizer (32K vocab, BPE)
- CLI for tokenization/detokenization
- Tokenizer validation (32K vs 50K trade-off analysis)
- Prototype training pipeline on Mac MPS
- Checkpointing and evaluation automation

Not Started:

- SLM model training pipeline (3B-7B parameters)
- LoRA fine-tuning setup
- RAG (Retrieval-Augmented Generation) integration
- Tool-using agent framework
- Physics paper dataset preparation at scale

9.3 Phase 1: Environment & Data Pipeline (Week 1-2)

Framework: Python 3.10+, PyTorch (with MPS), Hugging Face Transformers, Datasets, Accelerate, Pandas, NumPy

Steps:

1. Activate venv: `source /Users/jdsingh/slm_v0/venv/bin/activate`
2. Install packages: `pip install torch transformers datasets accelerate peft`
3. Verify MPS support for Mac
4. Download 10K-50K physics papers from arXiv
5. Organize into: `/Users/jdsingh/slm_v0/data/physics_papers/raw/`
6. Extract abstracts, introductions, methodologies
7. Remove noise, duplicates, non-English text
8. Normalize whitespace, split into chunks (max 4K words)
9. Tokenize using Rust tokenizer (32K vocab, frozen)

10. Split into train/val/test (98%/1%/1%), pack into 512-token sequences

****Expected Outcome:**** 300M-1B tokens ready for SLM training.

9.4 Phase 2: Base SLM Training (Week 3-5)

****Framework:**** PyTorch, Hugging Face Transformers, Accelerate

Parameter	3B Model	7B Model
Architecture	GPT2 or LLaMA-style decoder-only	GPT2 or LLaMA-style decoder-only
Hidden size	2048	3072
Layers	24	32
Attention heads	16	24
Max sequence length	512	512
Vocab size	32,000	32,000
Batch size	32	32
Gradient accumulation	4	4
Learning rate	5e-4 with cosine schedule	5e-4 with cosine schedule
Total steps	100,000	100,000

****Success Criteria:****

- Training loss decreases smoothly
- Validation perplexity < 20
- Generated samples coherent
- Peak memory < 16GB (Mac)

****Expected Outcome:**** Production-ready base SLM (3B-7B parameters).

9.5 Phase 3: LoRA Fine-tuning (Week 6-7)

****Framework:**** PEFT (Parameter-Efficient Fine-Tuning)

Parameter	Value
Rank	32
Alpha	32
Dropout	0.05
Target modules	q_proj, v_proj
Trainable parameters	~1-2% of base model

****Training:****

- Batch size: 16
- Learning rate: 2e-4
- Epochs: 3-5
- Output: LoRA weights (~50-100MB)

****Expected Outcome:**** Fine-tuned SLM ready for RAG integration.

9.6 Phase 4: RAG Integration (Week 8-9)

****Framework:**** FAISS, Sentence-Transformers, LangChain

****Steps:****

1. Generate embeddings for all physics paper chunks
2. Build FAISS vector index (384 or 768 dimensions)
3. Create retrieval function: query -> top-k relevant chunks
4. Augment prompt with retrieved context
5. Generate answer using fine-tuned SLM

****Expected Outcome:**** RAG system with physics-grounded responses.

9.7 Phase 5: Tool-Using Agents (Week 10-11)

****Framework:**** LangChain, LanGraph (ReAct framework)

Tool	Purpose
search_arxiv	Query arXiv for physics papers
solve_equation	Solve math equations (SymPy)
execute_python	Safe code execution for calculations

****Agent Loop:****

1. Generate thought + action
2. Execute appropriate tool
3. Observe result
4. Iterate until final answer
5. Max steps: 10-15

****Expected Outcome:**** Multi-step reasoning with tool integration.

9.8 Phase 6: Deployment (Week 12+)

****Framework:**** FastAPI, Docker

Endpoint	Method	Description
/query	POST	Submit physics question
/status	GET	Model health check
/upload_papers	POST	Add new physics papers to RAG

****UI (Optional):**** Streamlit web interface

****Deployment:**** Docker container on cloud (AWS, GCP, Azure, Hugging Face Spaces)

****Expected Outcome:**** Production-ready Physics Research Assistant.

9.9 Timeline

Week	Phase	Deliverable
1-2	Data Pipeline	300M-1B tokenized tokens
3-5	Base SLM	3B-7B model checkpoint
6-7	LoRA	50MB LoRA weights
8-9	RAG	Vector index + retrieval
10-11	Agents	Multi-step reasoning demo

| 12+ | Deployment | API + UI live |

10 Hardware Requirements

Level	Specs
Minimum (M1/M2 Mac)	16GB RAM, 100GB SSD
Recommended	GPU with 24GB+ VRAM, 32-64GB RAM, 500GB NVMe SSD
Training (AWS p4d.24xlarge)	8x A100, 320GB total VRAM, spot instances

10.1 Success Criteria

- Generates coherent physics-grounded answers
- Uses tools appropriately
- Cites retrieved papers
- Handles multi-step reasoning
- API response time < 5s
- Deployment-ready

---\n\n## File-Level Reference

- `subword\_tokenizer/src/lib.rs`: tokenizer core and model operations
- `subword\_tokenizer/src/main.rs`: CLI dispatch (`bpe-tokenizer`)
- `stream\_train.py`: prototype streaming training loop
- `scripts/evaluate\_checkpoints.py`: checkpoint ranking
- `scripts/run\_prototype\_3x\_and\_select\_best.sh`: train+select workflow
- `scripts/run\_prototype\_long\_4h.sh`: long prototype preset

Architecture Diagram

![System Architecture](architecture.png)

Figure 1: Live arXiv Stream → Rust Tokenizer → TinyLM Training → Checkpoint Evaluation

Live arXiv Stream → stream_train.py → Retry/Backoff → Rust Tokenizer CLI → Token IDs → TinyLM Training → Checkpoints → Evaluation → best_checkpoint_*.json

Prepared on July 3, 2026

Updated July 5, 2026 — Added references and architecture diagram for IIT Bombay server request for end-of-day status handoff.

References

[1] P. Jhandi, O. Kazi, S. Subramanian, N. Sendas. "Small Language Models for Efficient Agentic Tool Calling: Outperforming Large Models with Targeted Fine-tuning." arXiv preprint arXiv:2512.15943,

2025.

- [2] Y. Kang, et al. "Fine-tuning Small Language Models as Efficient Enterprise Search Relevance Labelers." arXiv preprint arXiv:2601.03211, 2026.
- [3] R. Sharma, M. Mehta. "Small Language Models for Agentic Systems: A Survey of Architectures, Capabilities, and Deployment Trade-offs." arXiv preprint arXiv:2510.03847, 2025.
- [4] Z. K. Chong, et al. "Compiling Deterministic Structure into SLM Harnesses." arXiv preprint arXiv:2604.17450, 2026.
- [5] ServiceNow, SLB. "Scaling AI Development with DGX Cloud: ServiceNow and SLB Production Deployments." Nvidia DGX Cloud Case Study, ZenML LLMops Database, 2025.
- [6] N. Patience. "Leveraging Small Language Models for Enterprise AI: Benefits, Use Cases, and IBM's Approach." The Futurum Group, in partnership with IBM, 2025.
- [7] M. Elfeki, R. Liu, C. Voegelé. "Return of the Encoder: Maximizing Parameter Efficiency for Small Language Models." arXiv preprint arXiv:2501.16273, 2025.
- [8] SACAI R. "Small Language Models for Enterprise Edge Deployment." Southern African Conference on Artificial Intelligence Research, 2025.
- [9] "Data-Centric Fine-Tuning of Small Language Models for Industrial Applications." IEEE Transactions on Industrial Informatics, 2025.
- [10] LinkedIn Engineering. "Production-Scale SLM Ranking: Optimizations for Inference." arXiv preprint arXiv:2510.22101, 2025.
- [11] A. Karpathy. "NanoGPT: The simplest, fastest repository for training medium-sized GPTs." GitHub, 2023.
- [12] R. Eldan, Y. Li. "TinyStories: How Small Can Language Models Be and Still Speak Coherently?" arXiv preprint arXiv:2305.07759, 2023.